
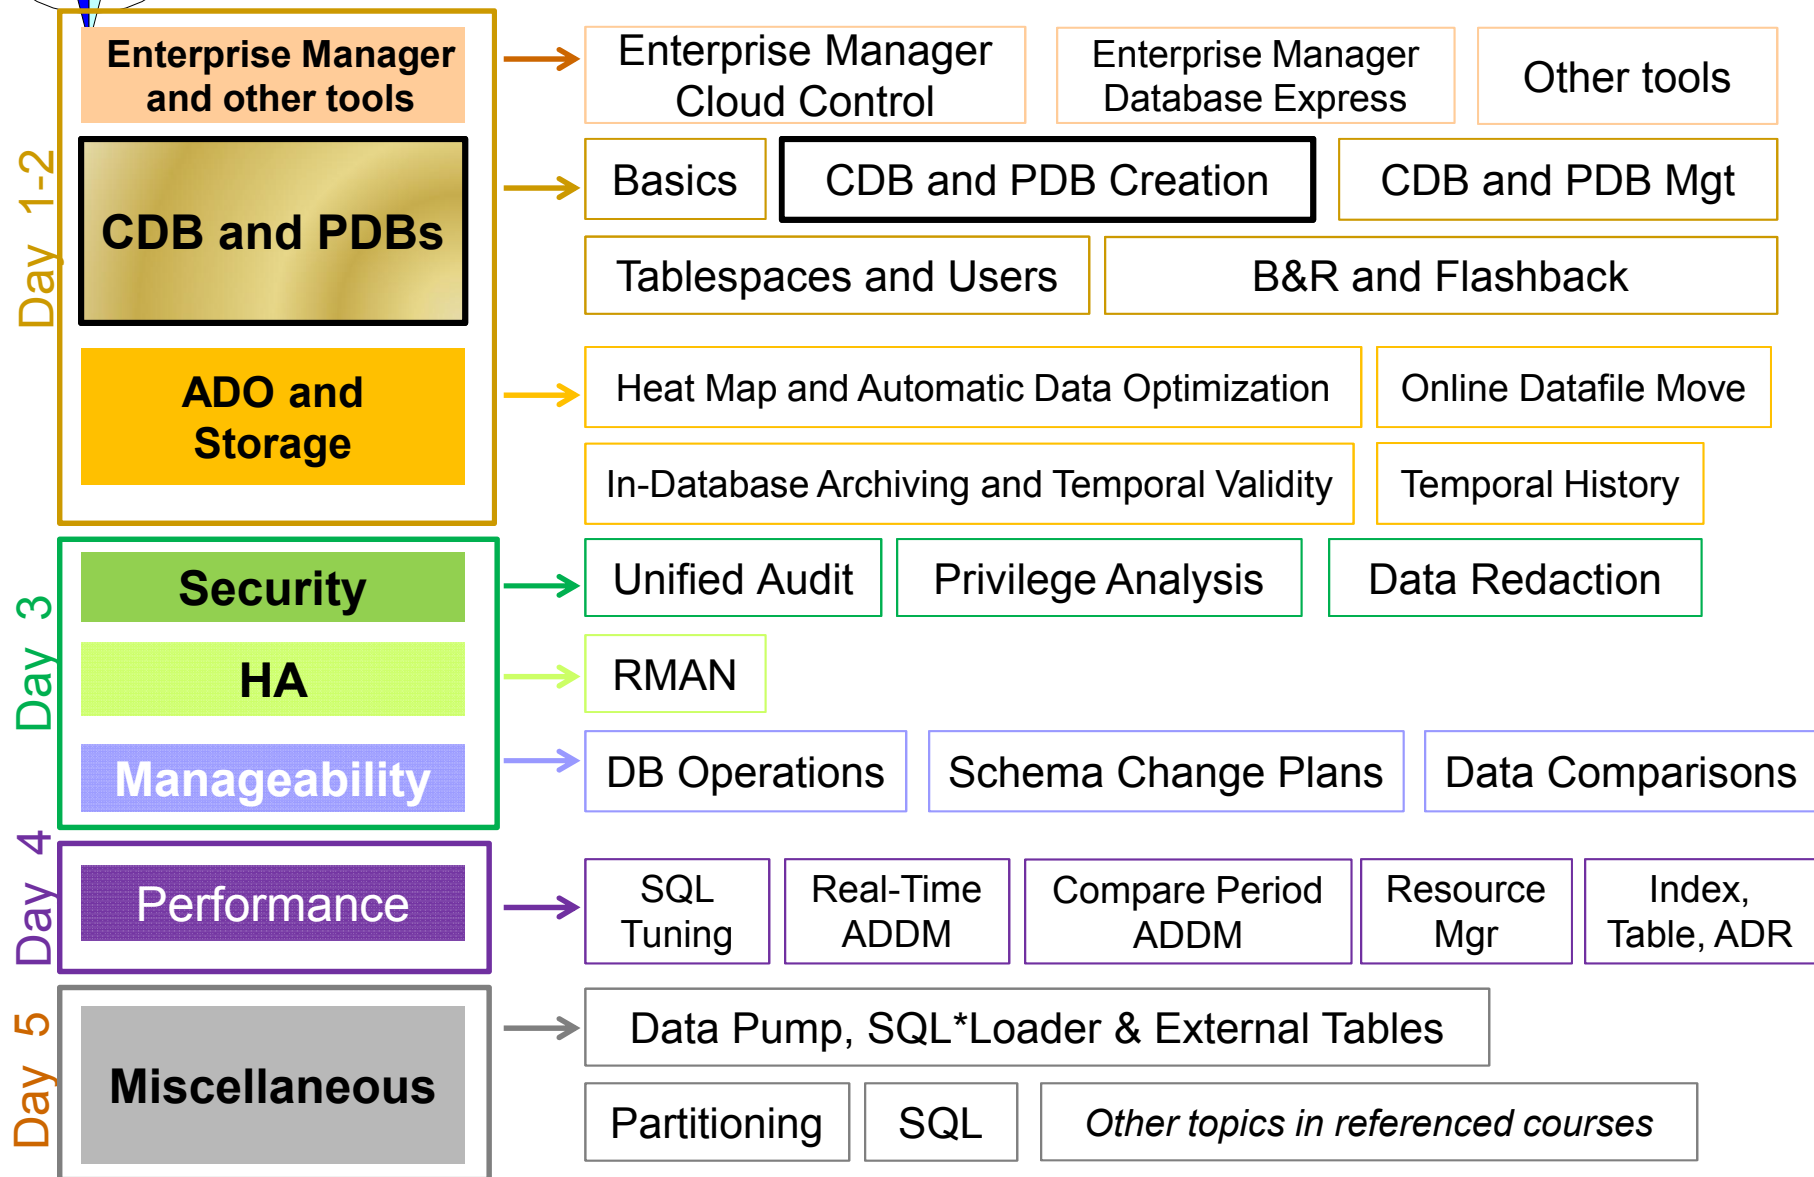



Creating Multitenant Container Databases and Pluggable Databases



Oracle Database 12c New and Enhanced Features






Objectives

After completing this lesson, you should be able to:


- Configure and create a CDB
- Create a PDB from `PDB$SEED`
- Create a PDB from a non-CDB
- Clone a PDB into the same CDB
- Unplug and plug a PDB from one CDB to another a CDB
- Explore the instance
- Explore the structure of PDBs
- Drop a PDB
- Migrate a pre-12.1 non-CDB database to CDB



Goals

Create a multitenant container database:

- To consolidate many pre-12.1 non-CDBs into a single, larger database
- To prepare a container
 - For plugging any future new application
 - For testing applications
 - For diagnosing application performance
- To simplify and reduce time for patching and upgrade

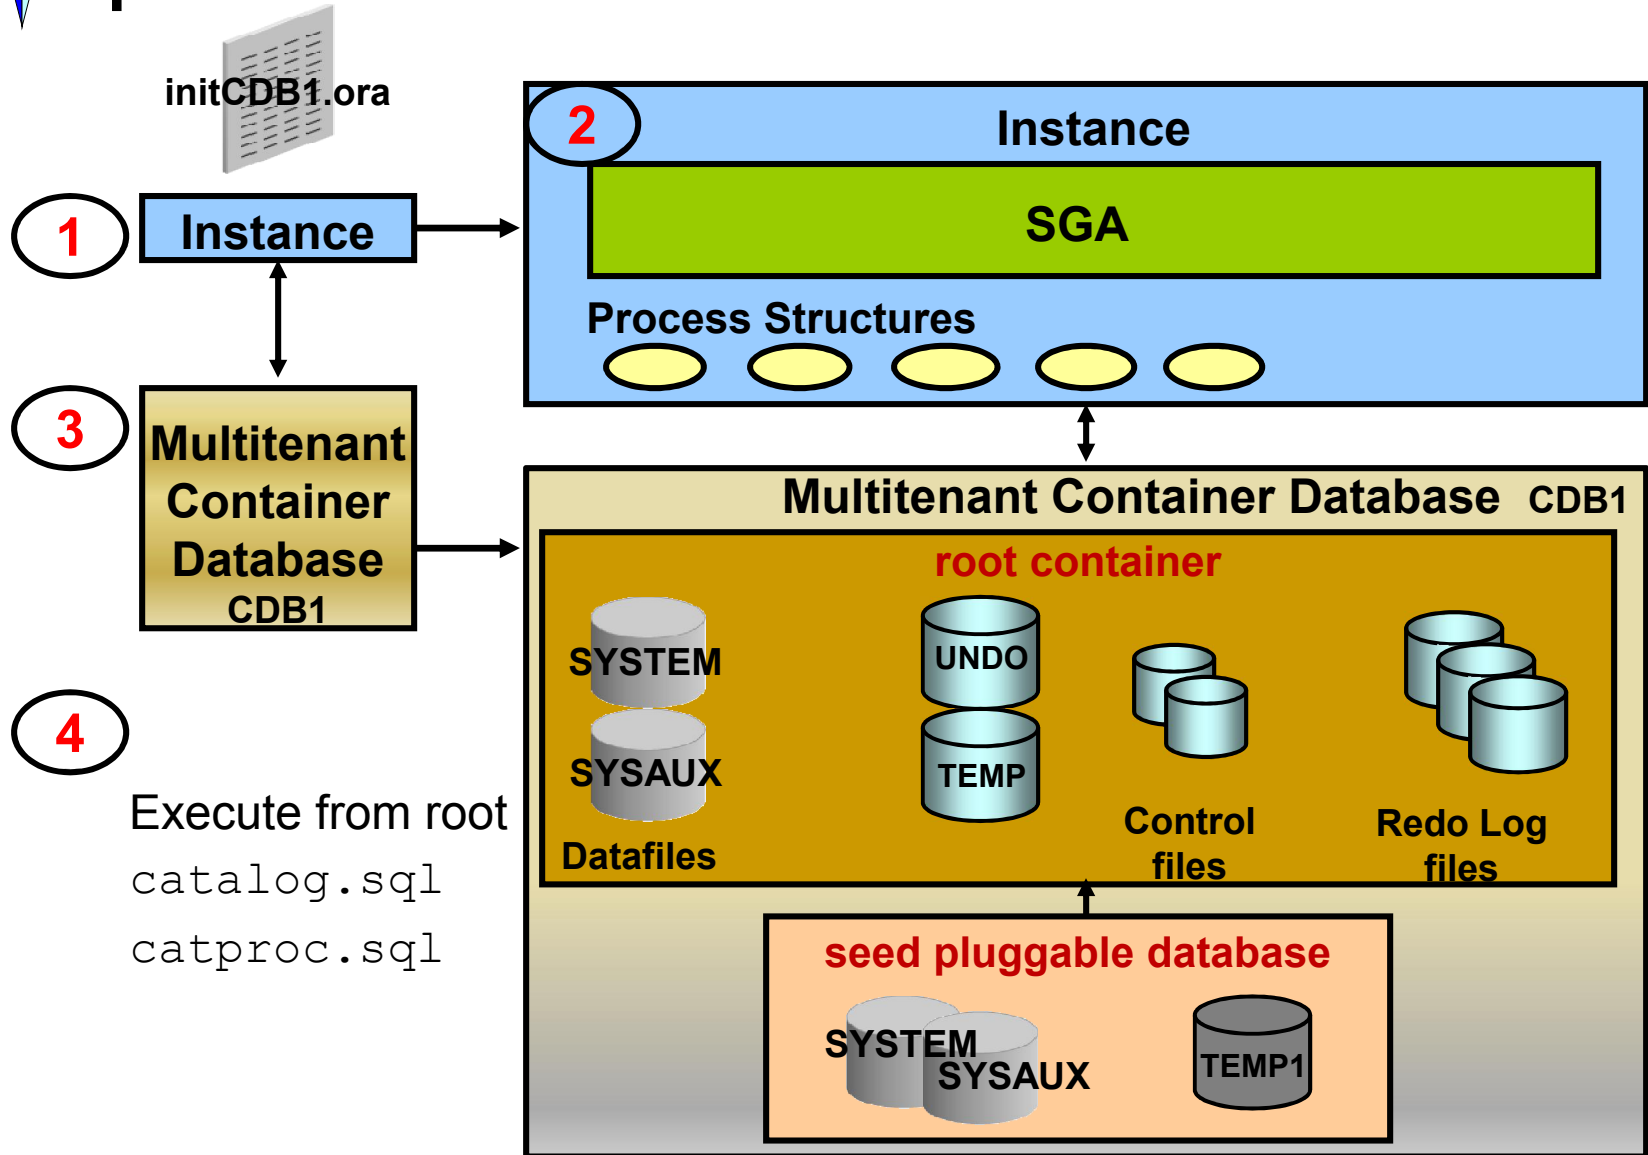



Tools

	SQL*Plus	OUI	DBCA	EM Cloud Control	SQL Developer	DBUA
Create a new CDB or PDB	Yes	Yes	Yes	Yes (PDB only)	Yes (PDB only)	
Explore CDB instance, architecture, PDBs	Yes			Yes	Yes	
Upgrade a 12.1 CDB to 12.x CDB	Yes			Yes		Yes



Steps to Create a Multitenant Container Database





Creating a Multitenant Container Database: Using SQL*Plus

1. Instance startup:

- a. Set `ORACLE_SID=CDB1`
- b. Set in `initCDB1.ora`:
 - Set `CONTROL_FILES` to CDB control file names
 - Set `DB_NAME` to CDB name
 - Set `ENABLE_PLUGGABLE_DATABASE` to `TRUE`


```
SQL> CONNECT / AS SYSDBA
SQL> STARTUP NOMOUNT
```


2. Create the database:

```
SQL> CREATE DATABASE CDB1 ENABLE PLUGGABLE DATABASE ...
      2 SEED FILE_NAME_CONVERT ('/oracle/dbs', '/oracle/seed');
```

- `CDB$ROOT` container
- `PDB$SEED` pluggable database

3. Close/open the seed PDB and run post-creation scripts.





Creating a Multitenant Container Database: Using DBCA

Database Configuration Assistant - Application - Step 2 of 5

Creation Mode

Database Operation
Creation Mode
Pre Requisite Checks
Summary
Progress Page

☒ Create a database with default configuration

Global Database Name: CDB1

Storage Type: File System

Database Files Location: {ORACLE_BASE}/oradata Browse...

Fast Recovery Area: {ORACLE_BASE}/fast_recovery_area Browse...

Administrative Password: ...

Confirm Password: ...


☒ Create As Container Database

Pluggable Database Name? PDB1

☐ Advanced Mode

Allows customization of storage locations, initialization parameters, management options, database options and different passwords for Administrator user accounts.


Help < Back Next > Finish Cancel



New Clause: SEED FILE_NAME_CONVERT

CREATE DATABASE new clauses:

```
SQL> CREATE DATABASE cdb1
 2  USER SYS IDENTIFIED BY p1 USER SYSTEM IDENTIFIED BY p2
 3  LOGFILE GROUP 1 ('/u01/app/oradata/CDB1/redo1a.log',
 4                  '/u02/app/oradata/CDB1/redo1b.log') SIZE 100M,
 5          GROUP 2 ('/u01/app/oradata/CDB1/redo2a.log',
 6                  '/u02/app/oradata/CDB1/redo2b.log') SIZE 100M
 7  CHARACTER SET AL32UTF8 NATIONAL CHARACTER SET AL16UTF16
 8  EXTENT MANAGEMENT LOCAL DATAFILE
 9          '/u01/app/oradata/CDB1/system01.dbf' SIZE 325M
10  SYSAUX DATAFILE '/u01/app/oradata/CDB1/sysaux01.dbf' SIZE 325M
11  DEFAULT TEMPORARY TABLESPACE tempts1
12          TEMPFILE '/u01/app/oradata/CDB1/temp01.dbf' SIZE 20M
13  UNDO TABLESPACE undotbs
14          DATAFILE '/u01/app/oradata/CDB1/undotbs01.dbf' SIZE 200M
15  ENABLE PLUGGABLE DATABASE
16  SEED FILE_NAME_CONVERT =
17          ('/u01/app/oradata/CDB1',
18          '/u01/app/oradata/CDB1/seed');
```



New Clause: ENABLE PLUGGABLE DATABASE

Without SEED FILE_NAME_CONVERT:

- OMF: DB_CREATE_FILE_DEST='/u01/app/oradata'
- Or new instance parameter:
PDB_FILE_NAME_CONVERT =
'/u01/app/oradata/CDB1','/u01/app/oradata/seed'

```
SQL> CONNECT / AS SYSDBA
```

```
SQL> STARTUP NOMOUNT
```

```
SQL> CREATE DATABASE cdb2
```

```
2   USER SYS IDENTIFIED BY p1 USER SYSTEM IDENTIFIED BY p2
```


```
3   EXTENT MANAGEMENT LOCAL
```

```
4   DEFAULT TEMPORARY TABLESPACE temp
```

```
5   UNDO TABLESPACE undotbs
```

```
6   DEFAULT TABLESPACE users
```


```
7   ENABLE PLUGGABLE DATABASE;
```



After CDB Creation: What's New in CDB

A CDB has new characteristics compared to non-CDBs:

- **Two containers:**
 - The **root** (CDB\$ROOT)
 - The **seed** PDB (PDB\$SEED)
- **Several services:** one per container
 - Name of root service = name of the CDB (cdb1)
- **Common** users in root and seed: SYS, SYSTEM ...
- **Common** privileges granted to common users
- **Pre-defined** common roles
- **Tablespaces** and data files associated to each container:
 - **root:**
 - SYSTEM: system-supplied metadata and no user data
 - SYSAUX
 - **seed:** SYSTEM, SYSAUX



Data Dictionary Views: DBA_XXX

DBA_XXX All of the objects in the root or a pluggable database

ALL_XXX Objects accessible by the current user in a PDB

USER_XXX Objects owned by the current user in a PDB

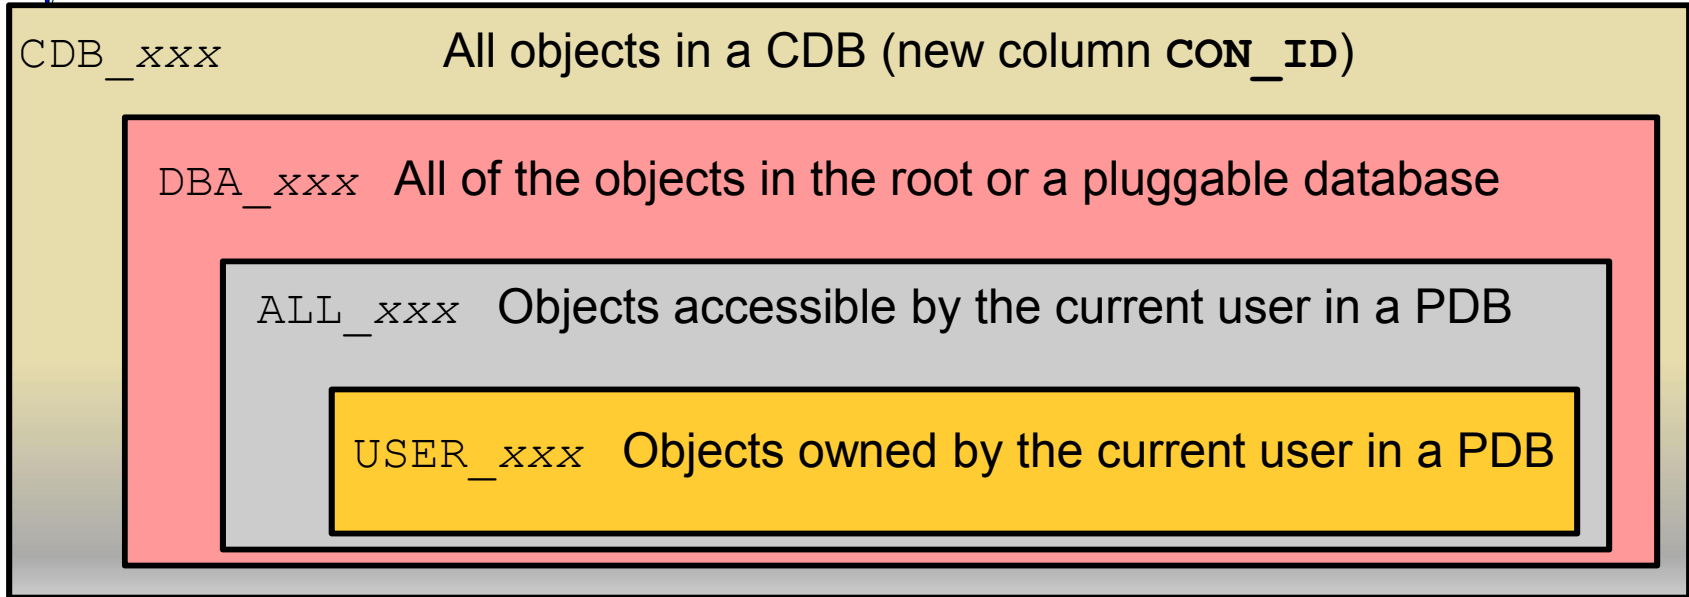
DBA dictionary views providing information within PDB:

```
SQL> SELECT table_name FROM dict
2 WHERE table_name like 'DBA%';
```

- **DBA_tablespaces:** All tablespaces of the PDB
- **DBA_data_files:** All data files of the PDB
- **DBA_tables:** All tables in the PDB
- **DBA_users:** All common and local users of the PDB




Data Dictionary Views: CDB_XXX



CDB dictionary views provide information across PDBs:

```
SQL> SELECT view_name FROM dba_views  
2 WHERE view_name like 'CDB%';
```

- CDB_pdb: All PDBS within the CDB
- CDB_tablespace: All tablespaces within the CDB
- CDB_data_files: All datafiles within the CDB
- CDB_users: All users within the CDB (common and local)



Data Dictionary Views: Examples

- Comparisons:

1

```
SQL> CONNECT / AS SYSDBA
SQL> SELECT role, common, con_id FROM cdb_roles;
```

2

```
SQL> SELECT role, common FROM dba_roles;
```

3

```
SQL> CONNECT sys@PDB1 AS SYSDBA
SQL> SELECT role, common, con_id FROM cdb_roles;
```

4


```
SQL> SELECT role, common FROM dba_roles;
```

- Access to data in V\$ or GV\$ views showing data from multiple PDBs can be secured using privilege.

```
SQL> SELECT name, open_mode FROM v$pdb;
```

NAME	OPEN_MODE

PDB\$SEED	READ ONLY
PDB1	READ WRITE
PDB2	READ WRITE



Data Dictionary Views: V\$xxx Views


SGA accessed by all containers: V\$ views and **CON_ID** column

```
SQL> SELECT distinct status, con_id FROM v$bh order by 2;
```

STATUS	CON_ID	
-----	-----	
cr	1	→ root
free	1	
xcur	1	
xcur	2	→ seed PDB
cr	3	→ PDB1 PDB
xcur	3	

```
SQL> select OBJECT_ID, ORACLE_USERNAME, LOCKED_MODE, CON_ID
2 from V$LOCKED_OBJECT;
```


OBJECT_ID	ORACLE_USERNAME	LOCKED_MODE	CON_ID	
-----	-----	-----	-----	
83711	SYS	3	3	← PDB1 PDB
83710	DOM	3	4	← PDB2 PDB



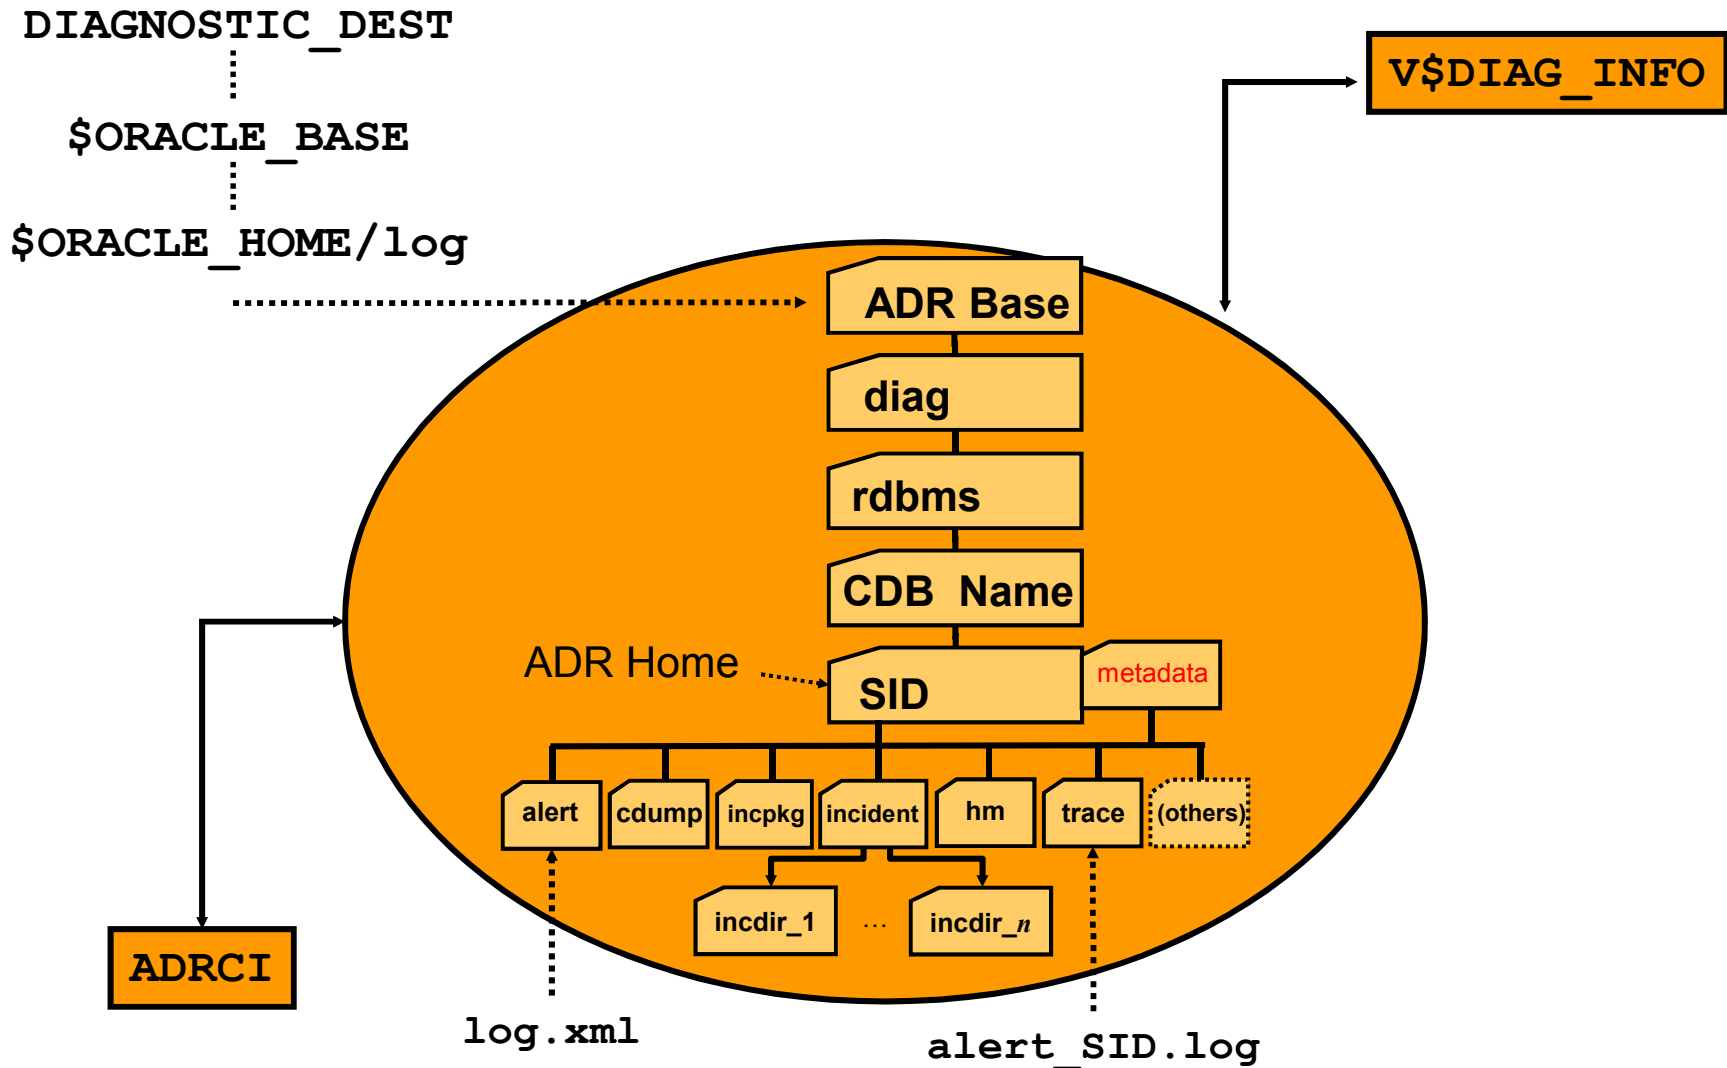
After CDB Creation: To-Do List


After CDB creation, the CDBA has to:

- Set a separate default tablespace for the root and for each PDB
- Set a default temporary tablespace for the entire CDB (optionally create additional temporary tablespaces in individual PDBs)
- Start the listener
- Plug non-CDBs
- Test startup/shutdown procedures
- Create new event triggers to automate PDBs opening
- Create backup and recovery procedures



Automatic Diagnostic Repository






Automatic Diagnostic Repository: alert.log File

The alert_CDB1.log shows new DDL statements.

```
CREATE DATABASE cdb1
...
ENABLE PLUGGABLE DATABASE
SEED
FILE_NAME_CONVERT=('/u01/app/oradata/CDB1','/u01/app/oradata
/seed');


CREATE PLUGGABLE DATABASE pdb1 ... ;
ALTER PLUGGABLE DATABASE pdb1 UNPLUG INTO ... ;
ALTER PLUGGABLE DATABASE ALL OPEN ;
ALTER PLUGGABLE DATABASE pdb2 CLOSE IMMEDIATE ;
```



Quiz

Which is a characteristic of the seed pluggable database of a CDB?


- a. It is always kept in `READ ONLY` mode.
- b. It is not a container.
- c. The seed can be dropped.



Quiz

You create a CDB. What is true about the seed pluggable database?


- a. Copy the seed data files yourself.
- b. Use the new clause `SEED FILE_NAME_CONVERT` in the `CREATE DATABASE` statement.
- c. The seed pluggable database is not required.
- d. The seed pluggable database does not require data files.



Practice 3 Overview: Creating a CDB and PDBs

The first practice covers the following topic:


- Creating a CDB with no PDBs



Provisioning New Pluggable Databases

Four methods:


- Create a new PDB from the seed PDB.
- Plug a non-CDB in a CDB.
- Clone a PDB from another PDB:
 - Into the same CDB
- Plug an unplugged PDB into another CDB.



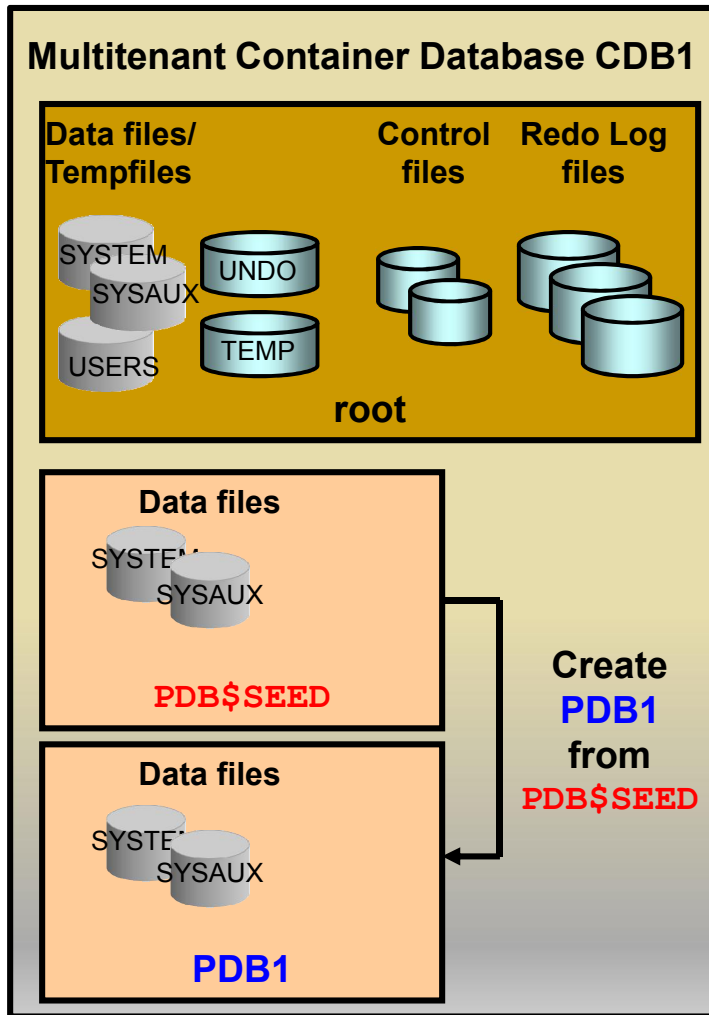
Tools

To provision new PDBs, you can use:


- SQL*Plus
- SQL Developer
- Enterprise Manager Cloud Control
- DBCA
 - Copy from seed
 - By unplugging / plugging method



Method 1: Create New PDB from PDB\$SEED



- Copies the data files from PDB\$SEED data files
- Creates tablespaces SYSTEM, SYSAUX
- Creates a full catalog including metadata pointing to Oracle-supplied objects
- Creates common users:
 - Superuser SYS
 - SYSTEM
- Creates a local user (PDBA) granted local PDB_DBA role
- Creates a new default service



Steps: With `FILE_NAME_CONVERT`


Create a new PDB from the seed using **`FILE_NAME_CONVERT`**:

1. Connect to the root as a common user with `CREATE PLUGGABLE DATABASE` system privilege:

```
SQL> CREATE PLUGGABLE DATABASE pdb1
      2 ADMIN USER admin1 IDENTIFIED BY p1 ROLES=(CONNECT)
      3 FILE_NAME_CONVERT = ('PDB$SEEDdir', 'PDB1dir');
```

2. Use views to verify:

```
SQL> CONNECT / AS SYSDBA
SQL> SELECT * FROM cdb_pdfs;
SQL> SELECT * FROM cdb_tablespace;
SQL> SELECT * FROM cdb_data_files;
SQL> ALTER PLUGGABLE DATABASE pdb1 OPEN RESTRICTED;
SQL> CONNECT sys@pdb1 AS SYSDBA
SQL> CONNECT admin1@pdb1
```



Steps: Without `FILE_NAME_CONVERT`

Create a new PDB from seed without `FILE_NAME_CONVERT`:


- OMF: `DB_CREATE_FILE_DEST =`
`' /u01/app/oradata/CDB1/pdb1 '`

```
SQL> CREATE PLUGGABLE DATABASE pdb1  
2 ADMIN USER pdb1_admin IDENTIFIED BY p1  
3 ROLES=(CONNECT) ;
```

Or

- New parameter: `PDB_FILE_NAME_CONVERT =`
`' /u01/app/oradata/CDB1/seed', '/u01/app/oradata/CDB1/pdb1 '`

```
SQL> CREATE PLUGGABLE DATABASE pdb1  
2 ADMIN USER pdb1_admin IDENTIFIED BY p1  
3 ROLES=(CONNECT) ;
```



Method 1: Using SQL Developer

1. Select Container Database.


2. Click Create Pluggable...

3. Provide Properties information.

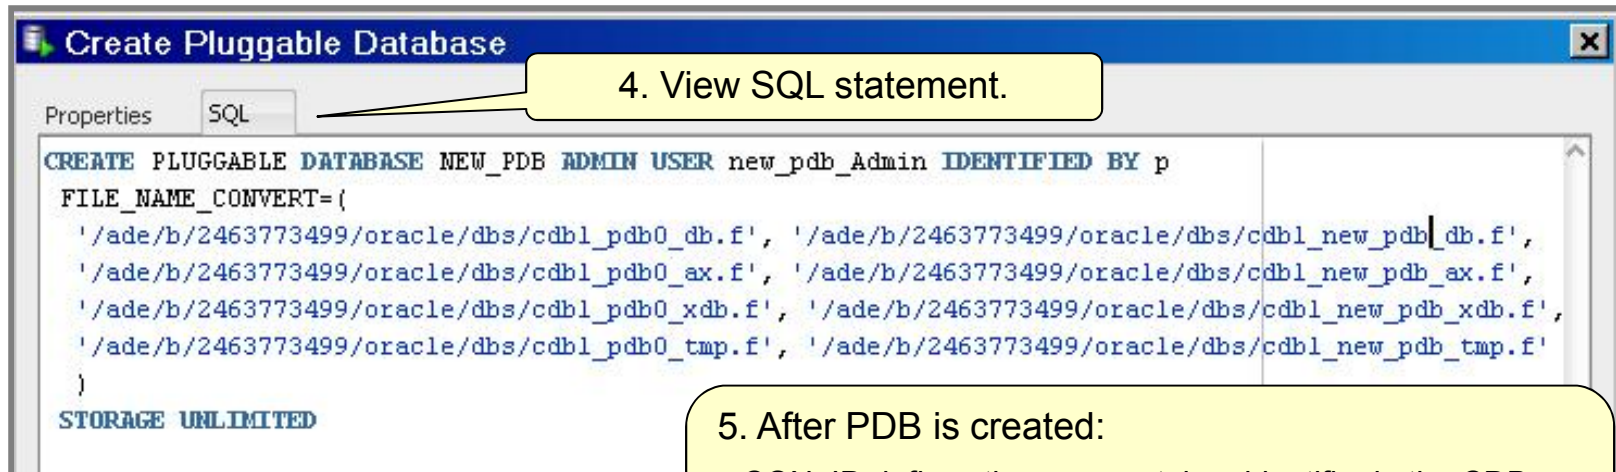
The screenshot shows the SQL Developer interface. The left pane displays the 'Connections' tree with 'CDB_ROOT_dom_as_sys' selected. The 'Container Database' folder is expanded, showing 'DOM_PDB4', 'DOM_PDB4I', 'NEW_PDB', 'PDB\$SEED', 'PDB1', 'PDB2', and 'PDB3'. The 'Create Pluggable...' option is highlighted in the context menu. The right pane shows the 'Create Pluggable Database' dialog box with the 'Properties' tab selected. The 'Database Name' is 'NEW_PDB', 'Admin Name' is 'new_pdb_Admin', and 'Admin Password' is masked. The 'File Name Conversions' section shows 'Custom Names' selected, with a table of source and target file paths. The 'Storage' section shows 'Unlimited' for 'Total Size' and 'Temp Tablespace Usage'.

Source File	Target File
/ade/b/2...	/ade/b/2463773499/oracle/dbs/cdb1_new_pdb_xdb.f
/ade/b/2...	/ade/b/2463773499/oracle/dbs/cdb1_new_pdb_tmp.f

Property	Value
Database Name	NEW_PDB
Admin Name	new_pdb_Admin
Admin Password	•
File Name Conversions	Custom Names
Storage	Unlimited
Total Size	Unlimited
Temp Tablespace Usage	Unlimited

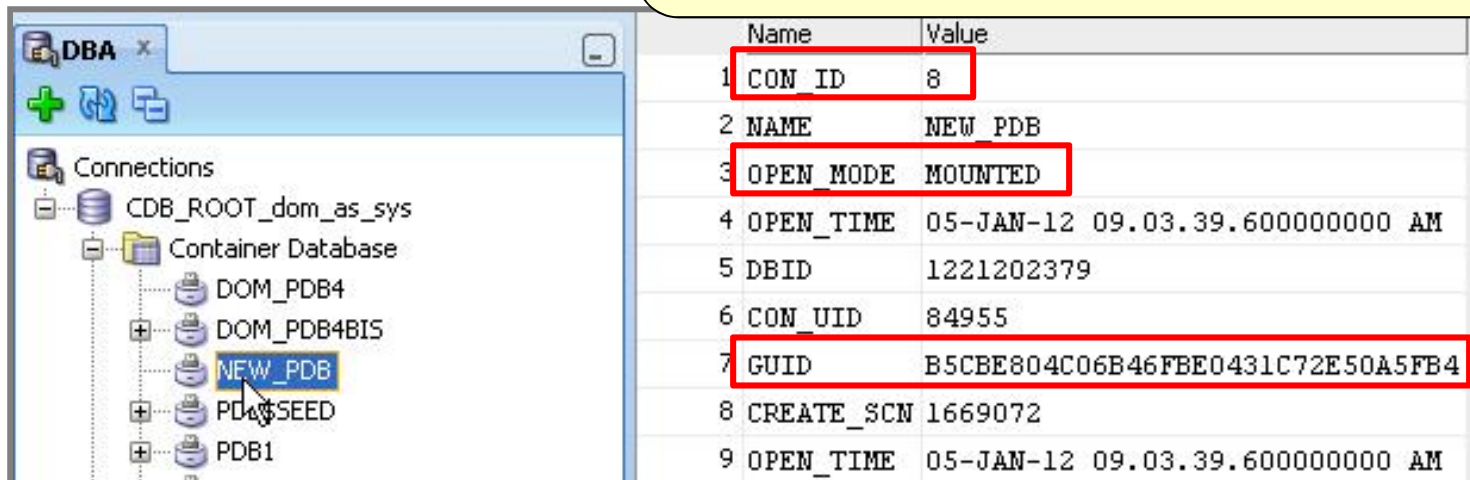



Method 1: Using SQL Developer



5. After PDB is created:


- CON_ID defines the new container identifier in the CDB.
- OPEN_MODE defines the open status, by default MOUNTED.
- GUID defines the new global unique container identifier.



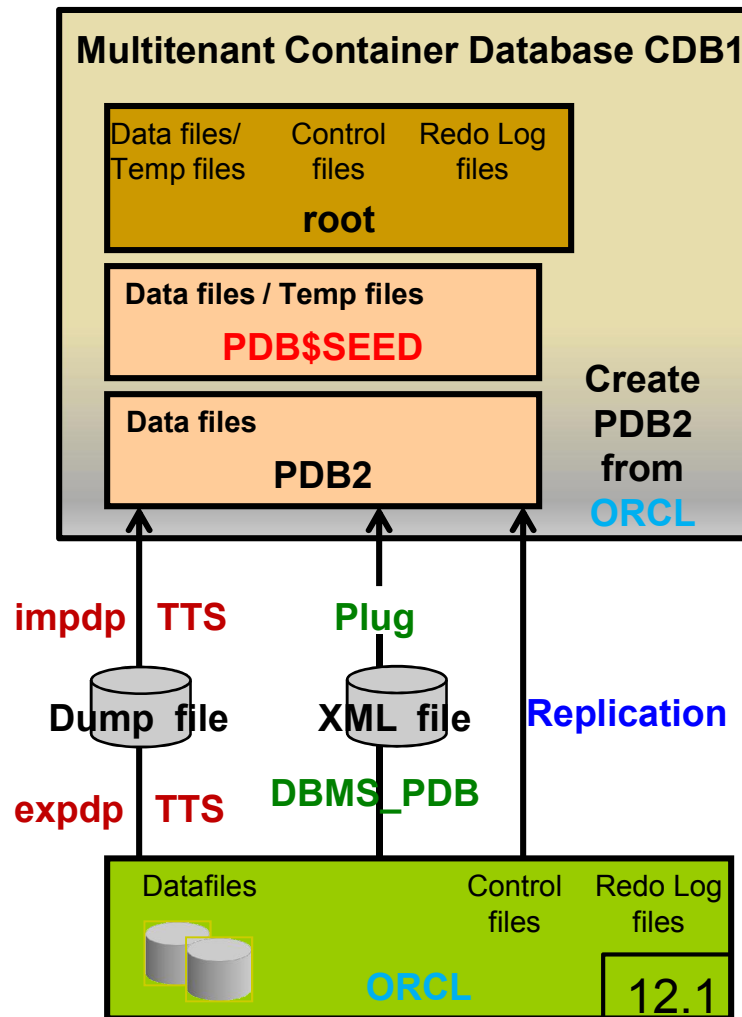


Synchronization

1. Customer-created common users or roles in root:
 - **Cannot** be created, modified, dropped when PDB is in `READ-ONLY` mode
 - Can be created, modified, dropped when PDB is in `MOUNTED` mode
2. When opening the PDB:
 - In `READ-ONLY` mode, an error is returned.
 - In `READ-WRITE` mode, synchronization with the target CDB is automatically completed.
 - A compatibility check is automatically performed:
 - Any violation is reported in the `PDB_PLUG_IN_VIOLATIONS` view.
 - If there are no violation, the PDB status is changed to `NORMAL`.



Method 2: Plug a Non-CDB into CDB




Three possible methods:

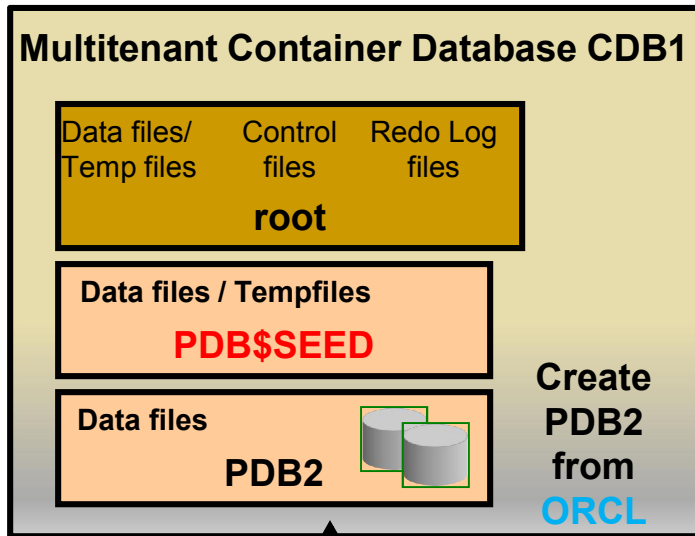
- **TTS or TDB or full export/import**
- XML file definition with `DBMS_PDB`
- **Replication**

Entities are created in the new PDB:

- Tablespaces: `SYSTEM`, `SYSAUX`
- A full catalog
- Common users: `SYS`, `SYSTEM`
- A local administrator (PDBA)
- A new default service




Plug a Non-CDB in to CDB Using DBMS_PDB



Plug

XML metadata file

DBMS_PDB.DESCRIBE




1. Open **ORCL** in READ ONLY mode
2.

```
SQL> EXEC DBMS_PDB.DESCRIBE  
(' /tmp/ORCL.xml ')
```
3. Connect to the target cdb as a common user with CREATE PLUGGABLE DATABASE privilege
4. Plug in the unplugged PDB **ORCL** as PDB2

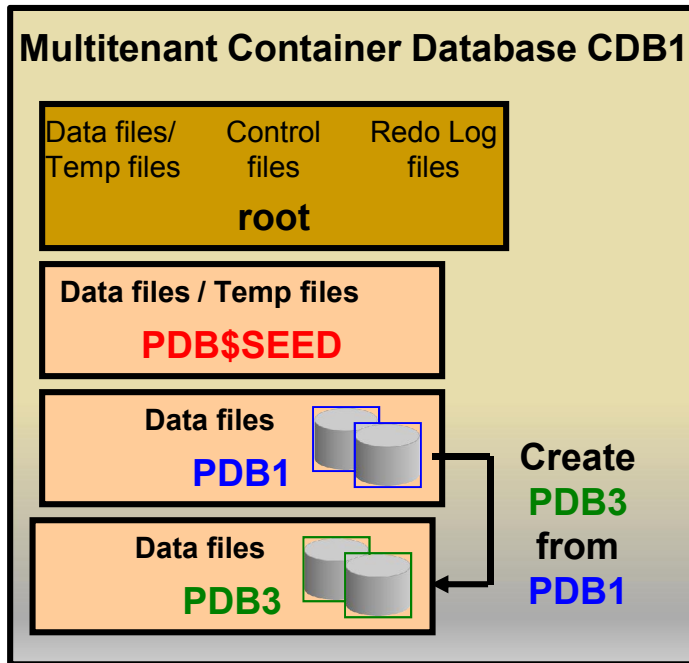
```
SQL> CREATE PLUGGABLE DATABASE  
2 PDB2 USING '/tmp/ORC.xml';
```
5. Run the noncdb_to_pdb.sql script

```
SQL> CONNECT sys@PDB2 AS SYSDBA  
SQL> @$ORACLE_HOME/rdbms/admin/noncdb_to_pdb
```
6. Open PDB2

Note: The STATUS of the PDB is CONVERTING.



Method 3: Clone PDBs



PDB3 owns:

- SYSTEM, SYSAUX tablespaces
- Full catalog
- SYS, SYSTEM common users:
- Same local administrator name
- New service name

1. In init.ora, set `DB_CREATE_FILE_DEST= 'PDB3dir'` or
`PDB_FILE_NAME_CONVERT='PDB1dir', 'PDB3dir'`

2. Connect to the root.


3. Quiesce **PDB1** (Close **PDB1** before):

```
SQL> ALTER PLUGGABLE DATABASE  
2   pdb1 OPEN READ ONLY;
```

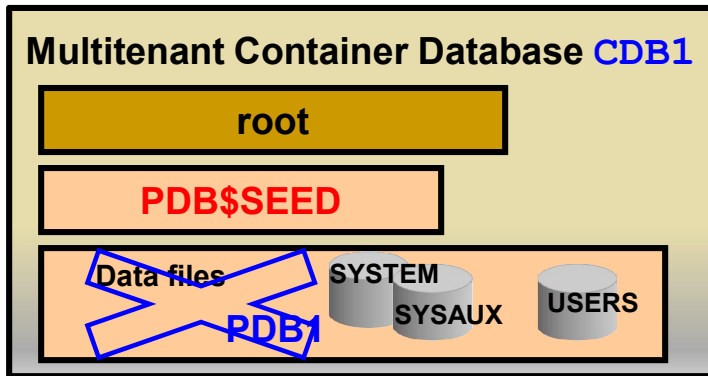
4. Clone **PDB3** from **PDB1**:

```
SQL> CREATE PLUGGABLE DATABASE  
2   pdb3 FROM pdb1;
```

5. Open **PDB3** in read-write mode.




Method 4: Plug Unplugged PDB in to CDB



Unplug PDB1



Plug PDB1



Unplug **PDB1** from **CDB1**:


1. Connect to **CDB1** as a common user.
2. Verify that **PDB1** is opened READ ONLY.
3.

```
SQL> ALTER PLUGGABLE DATABASE  
2   pdb1 UNPLUG INTO  
3       'xmlfile1.xml';
```
4. Drop **PDB1** from **CDB1**.

Plug **PDB1** in to **CDB2**:

1. Connect to **CDB2** as a common user.
2. Use `DBMS_PDB` package to check the compatibility of **PDB1** with **CDB2**.
3.

```
SQL> CREATE PLUGGABLE DATABASE  
2   pdb1 USING 'xmlfile1.xml'  
3   NOCOPY;
```
4. Open **PDB1** in read-write mode.



Method 4: Flow

Several clauses can be used in conjunction:

Are new PDB files based on same files that were used to create existing PDB in CDB?

If not, `AS CLONE` clause is required and so, it ensures that Oracle Database generates unique PDB DBID, GUID, and other identifiers expected for the new PDB.

XML file accurately describes current locations of files?

If not, the `SOURCE_FILE_NAME_CONVERT` clause is required.

Are files are in correct location?


If not, specify `COPY` to copy files to new location or `MOVE` to move them to another location.


If yes, use `NOCOPY`. `COPY` is the default.

- `FILE_NAME_CONVERT` clause of `CREATE PLUGGABLE DATABASE` statement
- `OMF: DB_CREATE_FILE_DEST` parameter
- `PDB_FILE_NAME_CONVERT` parameter

Do you want to specify storage limits for PDB?

If yes, specify the `STORAGE` clause.





Plug Sample Schemas PDB: Using DBCA

Database Configuration Assistant - Application - Step 1 of 7

Database Operation 1

Select the operation that you want to perform:

- ☒ **Create a Pluggable Database** 2
- ☐ Unplug a Pluggable Database
- ☐ Delete a Pluggable Database

Select the database in which Pluggable database needs to be created. 3

Select	Database
<input type="radio"/>	cdb1
<input type="radio"/>	cdb2
<input type="radio"/>	orcl
<input type="radio"/>	orcl3
<input checked="" type="radio"/>	cdb3

Database Configuration Assistant - Application - Step 4 of 7

Create Pluggable Database 4

ORACLE 12c DATABASE

Select the operation that you want to perform:


- ☐ Create a new Pluggable Database
- ☐ Create Pluggable Database From PDB Archive
- ☒ **Create Pluggable Database using PDB File Set**

Pluggable Database Archive:

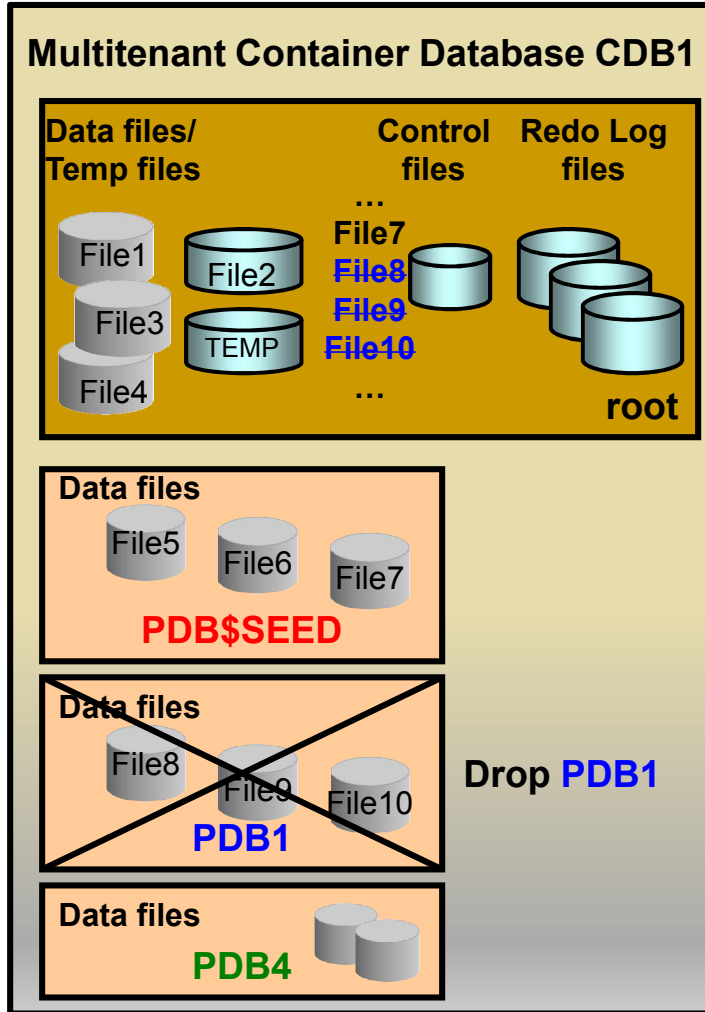
Pluggable Database Metadata File:

Pluggable Database Datafile Backup:

Plug a new PDB with Sample Schemas using a PDB File Set




Dropping a PDB



```
SQL> ALTER PLUGGABLE DATABASE
2   pdb1 CLOSE;
SQL> DROP PLUGGABLE DATABASE
2   pdb1 [INCLUDING DATAFILES] ;
```

- Updates control files
- If INCLUDING DATAFILES :
 - Removes PDB1 datafiles
- If KEEP DATAFILES (default):
 - Retain data files
 - Can be plugged in another or the same CDB
- Requires SYSDBA privilege
- Cannot drop seed PDB



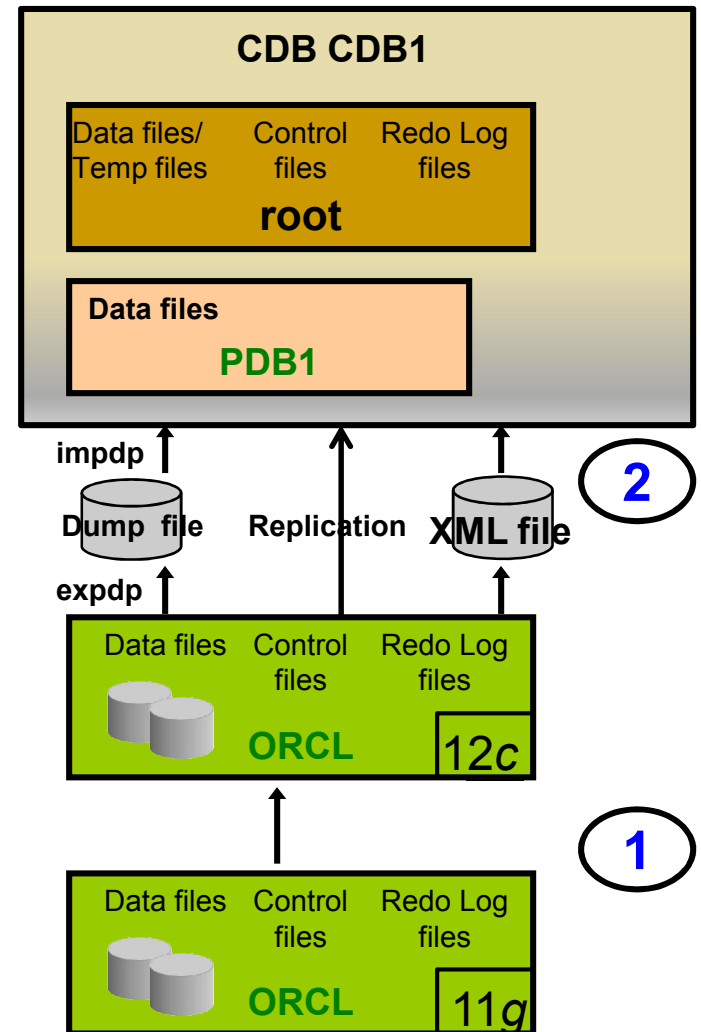
Migrating pre-12.1 Databases to 12.1 CDB


There are two methods:

1. Upgrade an existing pre-12.1 database to 12c.
2. Plug-in non-CDB into a CDB.

Or

1. Pre-create a PDB in CDB.
2. Use 11g expdp / 12c impdp or replication between non-CDB and PDB.






Quiz

Which of the following are true about cloning a PDB into the same CDB? Select all that apply.


- a. It is not possible. You can only clone a PDB into another CDB.
- b. You can clone only one PDB into the same CDB.
- c. Cloning a PDB can use the source files copy method to the target PDB files.
- d. Cloning a PDB can use the clause NOCOPY if the target PDB files will use the source files.



Quiz

Which of the following are true about dropping a PDB?


- a. You can drop a PDB only if the PDB is closed.
- b. You can possibly drop the seed PDB, but you will not be able to create any other PDB within the CDB.
- c. You can drop a PDB and keep the data files to be reused by another PDB.
- d. When you drop a PDB, the data files and redo log files are automatically removed from the storage file system.



Summary

In this lesson, you should have learned how to:

- Configure and create a CDB
- Create a PDB from `PDB$SEED`
- Create a PDB from a non-CDB
- Clone a PDB into the same CDB
- Unplug and plug a PDB from one CDB to another a CDB
- Explore the instance and structure of PDBs
- Drop a PDB
- Migrate pre-12.1 non-CDB database to CDB



Practice 3 Overview: Creating a CDB and PDBs

These practices cover the following topics:

- Creating a new PDB into a CDB using the seed
- Cloning a PDB from a CDB into the same CDB
- Plugging a non-CDB in to a CDB
- Merging two CDBs into a single one
- Dropping a PDB