

CHAPTER 10

TOPIC OBJECTIVES



- **Using Explain Plan**
- **Optimizer Rules**
- **Tuning SQL Statements**

EXPLAIN PLAN




What is going on with my SQL Statements?

The EXPLAIN PLAN diagnostic tool allows you to observe how ORACLE executes SQL statements. The EXPLAIN PLAN statement displays the execution plan selected by the OPTIMIZER for SELECT, INSERT, UPDATE, and DELETE statements. Before you can use this performance diagnostic tool you must:

- Create a plan table
- Understand what information the optimizer will load into this table
- Be able to maintain and synchronize the rows in the plan table

THE PLAN_TABLE COLUMNS(2)




Three new columns to use

There are new columns in the PLAN-TABLE that help estimate resource cost of statements.

- CPU_COST - This column holds numbers that represent the CPU cost of the statement. This column will be NULL if the rule-based optimizer is used.
- IO_COST - This column holds numbers that represent the disk I/O of the statement. This column will be NULL if the rule-based optimizer is used.
- TEMP_SPACE - This column holds numbers that represent the size, in bytes, of any temporary space the statements might require. If the rule-based optimizer is used or if the statement does not require any temporary space, this column will be NULL.

To use these columns make sure you have generated statistics on the schema objects you are evaluating.

EXPLAIN PLAN SYNTAX



When using EXPLAIN PLAN, the actual SQL statement is not executed. Instead, a plan tree is inserted into the PLAN_TABLE. A query must then be written against the PLAN_TABLE to interpret the results.

The identification and informational columns of the PLAN_TABLE are:

COLUMN	DESCRIPTION
STMT DESCRIPTION TABLE_NAME	A description of the statement The table in which the output from EXPLAIN PLAN is stored. The default is PLAN_TABLE.
STATEMENT_ID	The statement description
TIMESTAMP	The date and time EXPLAIN PLAN was executed.
REMARKS	Comments about the EXPLAIN results.
OPERATION	The operation performed in this step.
OPTIONS	Options about the operation.
OBJECT_NODE	The database link used to reference an object.
OBJECT_OWNER	The user that owns the table or view
OBJECT_NAME	Database, table, or index name
OBJECT_INSTANCE	The position of the object as it occurred in the statement. Numbering is from left to right, outer to inner. View expansion results in unpredictable numbers.
OBJECT_TYPE	Description of the object
SEARCH_COLUMNS	Not currently used.
ID	The user assigned id number for this statement and run
PARENT_ID	The id of the next execution step that operates on the output of the ID step
POSITION	Order of processing for steps have the same PARENT_ID
COST	Places a COST NUMBER in this column if the cost-based optimizer was used.
OTHER	Other information for this run, such as the query for remote node execution.

OPERATIONS IDENTIFIED BY THE EXPLAIN PLAN

Operations column of the PLAN_TABLE	Description
SELECT STATEMENT UPDATE STATEMENT DELETE STATEMENT INSERT STATEMENT AND-EQUAL	Used for WHERE clause equalities. Each comparison includes a non-unique indexed column from which ROWIDs are obtained and intersected. Indexes are "MERGED".
CONNECT BY CONCATENATION COUNT FILTER FIRST ROW FOR UPDATE INDEX	Retrieval based on a tree-walk UNION ALL for a group of tables (ORs rule) A count operation A restriction of the rows returned A retrieval of only the first row of a query A retrieval that row locks on selected rows A retrieval from an index Options: UNIQUE SCAN index search for unique values RANGE SCAN index search for range between
INTERSECTION	A retrieval of rows common to two tables. The tables are sorted first.
MERGE JOIN	A join formed by merging two sorted sets of data Options: OUTER -- an outer join MINUS -- a retrieval of rows in table 1 but not in table 2
NESTED LOOPS	A join on two child operations (each row returned by the first child operation is then operated on by the second child operation). One of three methods used to join tables. The other two are SORT MERGE and INDEX CLUSTER
PROJECTION	Selecting a subset of the columns
REMOTE	Retrieval across SQL*NET
SEQUENCE	An operation involving the sequence generator
SORT	A retrieval of rows ordered by one or more columns
SORT UNIQUE	UNIQUE -- sort to produce unique values
SORT GROUP BY	GROUP BY -- sort for grouping
SORT JOIN	JOIN -- sort for merge join
SORT ORDER BY	ORDER BY -- sort for order by

EXPLAIN PLAN TABLE COLUMNS

Operations Column of the PLAN_TABLE	Description
TABLE ACCESS	A retrieval from a base table
TABLE ACCESS BY ROWID	BY ROWID table access by ROWID
TABLE ACCESS FULL	FULL able access by Full Table Scan
UNION	CLUSTER table access by Cluster key A retrieval of unique rows from two tables UNIONed together dropping the duplicates
VIEW	A retrieval from a view of a table
OPTIMIZER	Choose, RULE, FIRST ROWS, ALL_ROWS

In order to load the PLAN_TABLE with rows concerning a particular SQL statement you would execute the following SQL: Review and execute the scripts big2.sql and big3.sql in order to run the following explain plan selects. ORACLE> @big2 @big3

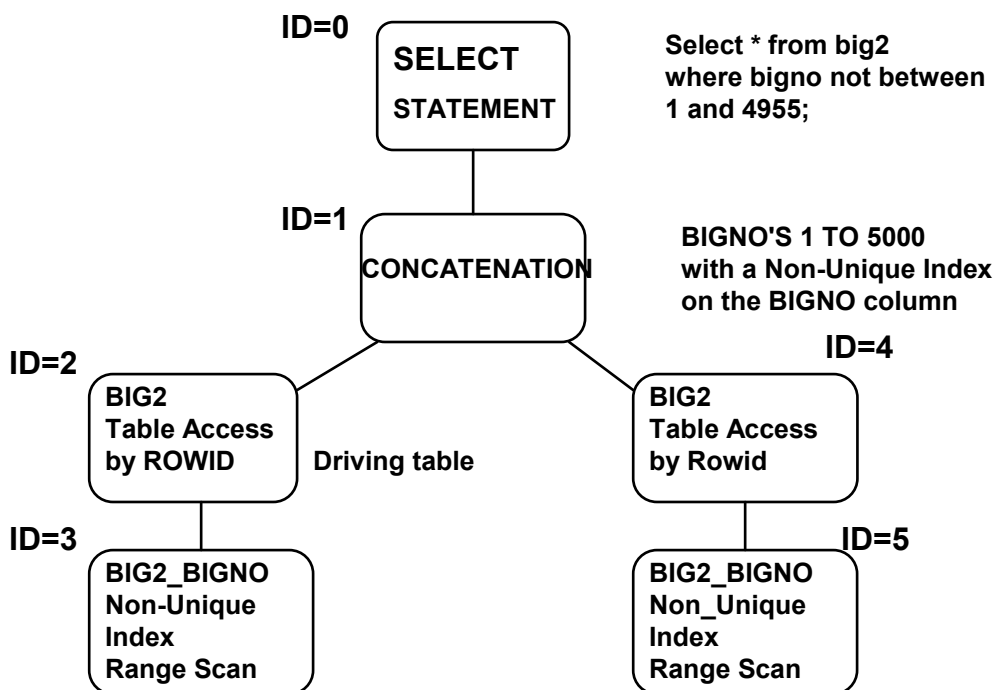
```
EXPLAIN PLAN
  [SET STATEMENT_ID = 'STMT DESCRIPTION']
  [INTO table_name]
  FOR SQL_STATEMENT;
```

EXAMPLES:

```
SQL> EXPLAIN PLAN FOR
      SELECT *
      FROM BIG2
      WHERE BIGNO NOT BETWEEN 1 AND 4955;
```

To see the plan results you would execute dbms_xplan(display)

PLAN_TREE ANALYSIS




select id, parent_id,object_name,object_type,operation,options
from plan_table order by id; OR
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY)

COL	OBJECT_NAME	FORMAT A12
COL	OBJECT_TYPE	FORMAT A12
COL	OPERATION	FORMAT A12
COL	OPTIONS	FORMAT A12

ID	PARENT_ID	OBJECT_NAME	OBJECT_TYPE	OPERATION	OPTIONS
0				Select statement	
1	0			Concatenation	
2	1	BIG2		Table Access	By ROWID
3	2	BIG2_BIGNO	Non-Unique	Index	RangeScan
4	1	BIG2		Table Access	By ROWID
5	4	BIG2_BIGNO	Non-Unique	Index	RangeScan

NESTED OUTPUT USING “START WITH” AND “CONNECT BY”



ORACLE> COL “Query Plan” format a80

```
ORACLE> select  LPAD ( ' ', 2*(LEVEL-1)) ||  
                OPERATION || ' ' ||  
                OPTIONS   || ' ' ||  
                OBJECT_NAME || ' ' ||  
                DECODE (ID, 0, 'COST = ' || POSITION)  
                “ Query Plan”  
from PLAN_TABLE  
START WITH ID = 0  
CONNECT BY PRIOR ID = PARENT_ID;
```

See explain.sql or the dbms_xplan command seeing the above code.
Query Plan

```
Select Statement          COST=1  
  Concatenation  
    Table access by ROWID BIG2  
      Index Range Scan Big2_BIGNO  
    Table access by ROWID BIG2  
      Index Range Scan BIG2_BIGNO
```

Read the output from the Explain Plan query by reading from the inside out and top bottom in a clockwise fashion.

- First, read the most indented clauses. If there is a tie, read from the top to the bottom

EXPLAIN PLAN OPERATION DEFINITIONS

Operation	Option	Definition
TABLE ACCESS		
	FULL	FULL TABLE SCAN - one row is read at a time until all rows are read.
	CLUSTER	Table scan on an Index Cluster Table Key
	HASH	A Hash Key is matched against a matching Hash value
	BY ROWID	Accesses a row by its ROWID usually from an Index
	UNIQUE SCAN	A single row is returned from a unique index
	FULL	A full scan of the index is performed. Because the index is in sorted order. It should be better than a full table scan and sort.
	FAST FULL	The rows do not have to be returned in sorted order. Utilizes multiblock I/O and can be run in parallel. All the blocks of the index are scanned
	RANGE SCAN	A range of values is returned from an index. Used with operators like <>, LIKE% or __, or BETWEEN.
	RANGE SCAN desc	A range scan in descending order
	AND EQUAL	Up to fine indexes can be merged. Each index is processed with the results merged with the results of another index.
UNION ALL		A set operation used for UNION ALL. Returns all rows from both queries. Does not eliminate duplicates or sort.
VIEW		A VIEW is a SET operation. A view will be evaluated and processed first, then merged with the rest of the query
HASH JOIN		One of the tables has in memory a bitmap created and a hashing function is used to find the matching rows in the second table.
MERGE JOIN		Tables are sorted and then merged
NESTED LOOPS		One of the tables is read one row at a time and a match is produced by accessing the second table. Usually, the second table has an index.
OUTER JOIN		With an OUTER JOIN, one of the tables is a complete table. All rows from that table will be returned regardless if there is a match. Performs a FULL TABLE SCAN.
SEQUENCE		when the NEXTVAL or CURRVAL is processed against a sequence number.

EXPLAIN PLAN OPERATIONS DEFINITIONS - PART 2

Operation	Option	Definition
SORT		Used with aggregate functions such as SUM, MIN, MAX, COUNT, and AVG
AGGREGATE		Used to sort rows with the ORDER BY clause
SORT ORDER BY		A row operation that will combine the results of an IN clause.
CONCATENATION		places a lock on the rows covered by the Select statement. This prevents updates against these rows while the query is running.
FOR UPDATE		restricts rows by a where clause i.e. WHERE EMPNO = 7900
FILTER		A set operation used for INTERSECT Returns rows common to both queries.
INTERSECTION		A set operation used for MINUS. Returns rows that are unique to the first query.
MINUS		A set operation used for UNION. Returns all rows that are unique to both queries. Performs a sort.
UNION		

EXPLAIN PLAN EXAMPLES

The following examples shows the plan for a two-table join where both sides of the WHERE clause are indexed:

```
DELETE FROM PLAN_TABLE;
ORACLE> alter session set OPTIMIZER_MODE = all_rows;
          (In this example choose and rule produce the same results)
ORACLE> EXPLAIN PLAN
        FOR
        SELECT BIG.BIGNO, BIG2.BNAME
        FROM BIG,
        BIG2
        WHERE BIG.BIGNO = BIG2.BIGNO;
```

```
ORACLE> COL "Query Plan" format a80
```

```
ORACLE> SELECT * FROM TABLE (dbms_xplan.display);
Query Plan
```

```
-----
Select Statement          COST=23
      NESTED LOOPS
            Table access   FULL BIG2
            Index Range Scan BIG_BIGNO
```

OPERATION	OBJECT_NAME	CARDINALITY	COST
SELECT STATEMENT		5000	1781
HASH JOIN		5000	1781
Access Predicates			
BIG.BIGNO=BIG2.BIGNO			
NESTED LOOPS		5000	1781
STATISTICS COLLECTOR			
INDEX (FAST FULL SCAN)	BIG2_BNAME_BIGNO	5000	7
INDEX (RANGE SCAN)	BIG_BIGNO	1	1741
Access Predicates			
BIG.BIGNO=BIG2.BIGNO			
INDEX (FAST FULL SCAN)	BIG_BIGNO	3000000	1741

OPTIMIZING SQL STATEMENTS

NOT TRANSLATIONS	
NOT STATEMENT	TRANSLATION
NOT >	<=
NOT >=	<
NOT <	>=
NOT <=	>
NOT BETWEEN	> larger number or < smallest number

The above chart indicates the replacements which can be used instead of the NOT statement. Some reasons for avoiding the NOT statement are:

- The use of !=, NOT=, <>, or NOT LIKE causes ORACLE not to use an index for that portion of the query
- The optimizer assumes that queries containing NOT usually retrieve more rows than they skip. So ORACLE will avoid using an Index.
- When NOT is used with mathematical operators other than =, ORACLE transforms it from a NOT so that indexes may be used.

The set operators UNION, MINUS and INTERSECT can usually be used to avoid NOT =.

NOTE: The query path is a result of weighing all search criteria paths against each other using the cost and determining the table which will be the driving table. Every AND is considered separately.

OPTIMIZING NOT BETWEEN

ORACLE> ALTER SESSION SET OPTIMIZER_MODE = ALL_ROWS;

(In this example FIRST_ROWS or ALL_ROWS PRODUCE THE SAME RESULTS)

```
ORACLE> EXPLAIN PLAN
        set statement_id = 'abc'
        FOR
        SELECT BNAME /* INDEX ON BIGNO */
        FROM BIG /* 2% OF THE ROWS RETURNED */
        WHERE BIGNO NOT BETWEEN 1000 AND 2999000;
```

ORACLE> COL "Query Plan" format a80

```
ORACLE> SELECT * FROM
TABLE(dbms_xplan.display('PLAN_TABLE','abc'))
ORACLE> SAVE explain
```

Query Plan

```
-----
Select Statement          COST=161
      CONCATENATION
            Table access BY ROWID BIG
                  Index Range Scan BIG_BIGNO
            Table access BY ROWID BIG
                  Index Range Scan BIG_BIGNO
```

Or

OPERATION	OBJECT_NAME	CARDINALITY	COST
SELECT STATEMENT		1999	24
VIEW	VW_ORE_A845E027	1999	24
UNION-ALL			
TABLE ACCESS (BY INDEX ROWID BATCH)		999	12
INDEX (RANGE SCAN)	BIG_BIGNO	999	5
Access Predicates			
BIGNO < 1000			
TABLE ACCESS (BY INDEX ROWID BATCH)		1000	12
INDEX (RANGE SCAN)	BIG_BIGNO	1000	5
Access Predicates			
BIGNO > 2999000			
Filter Predicates			
LNNVL(BIGNO < 1000)			

OPTIMIZING NOT EQUAL TO

```
ORACLE> ALTER SESSION SET OPTIMIZER_MODE = ALL_ROWS;
          (The use of FIRST_ROWS OR FIRST_ROWS_N produce the same results)
```

```
ORACLE> EXPLAIN PLAN
        set statement_id='abc2'
        FOR
        SELECT BNAME /* INDEX ON BIGNO */
        FROM BIG
        WHERE BIGNO <> 1;
```

```
ORACLE> @explain or
```


```
SELECT * FROM
```

```
TABLE (dbms_xplan.display('PLAN_TABLE','abc2'))
```

Query Plan

Select Statement	COST=317
TABLE ACCESS	FULL BIG

or



OPERATION	OBJECT_NAME	CARDINALITY	COST
SELECT STATEMENT			2401850
INDEX (FAST FULL SCAN)	BIG_BNAME_BIGNO		2467
Filter Predicates			
BIGNO <> 1			

ELIMINATING SQL STATEMENTS



Why is my SQL code running so long?

Use a single SQL statement rather than multiple statements whenever possible.

BAD:

```
UPDATE emp  
set ename = 'MOE'  
where empno = 3;
```

```
UPDATE emp  
set sal = 4000  
where empno = 3;
```

GOOD:

```
UPDATE emp  
set ename = 'MOE', SAL = 4000  
where empno = 3;
```


BAD:

```
SELECT * FROM EMP WHERE DEPTNO = 10;  
SELECT * FROM EMP WHERE DEPTNO = 20;  
SELECT * FROM EMP WHERE DEPTNO = 30;
```

GOOD:

```
SELECT * FROM EMP WHERE DEPTNO IN (10,20,30);
```

USING HINTS TO CONTROL THE DRIVING TABLE



Listed below are two methods of sending hints to the ORACLE optimizer:

Method 1: (Recommended)

```
ORACLE>  SELECT /*+ ORDERED */ ENAME,LOC
          FROM EMP, DEPT
          WHERE EMP.DEPTNO = DEPT.DEPTNO;
/* EMP IS THE DRIVING TABLE */
```

Method 2:

```
ORACLE>  SELECT /*+ USE_NL(DEPT,EMP) */ ENAME, LOC
          FROM DEPT, EMP
          WHERE EMP.DEPTNO = DEPT.DEPTNO;
-- DEPT IS THE DRIVING TABLE
```

HINTS:

- | | | |
|-------------|--------------|---------------|
| • FULL | • ROWID | • CLUSTER |
| • HASH | • INDEX | • INDEX |
| • INDEX_ASC | • INDEX_DESC | • INDEX_MERGE |
| • ORDERED | • USE_MERGE | • USE_NL |

Lab 9- using the optimizer and explain plan



Create the BIG and BIG2 tables and indexes using the script big1.sql and big2.sql Scripts. Given a non-unique index BIG_BIGNO on the BIGNO column of BIG, and a non-unique index BIG_BNAME on the BNAME column of the BIG table, use EXPLAIN PLAN to determine how ORACLE would process the following utilizing ALL_ROWS and FIRST_ROWS.

EXPLAIN PLAN FOR

1. ORACLE> Select * from big where bigno > 1;
SELECT * FROM TABLE (dbms_xplan.display

HINT: use AUTOTRACE to help provide the information.

2. ORACLE> delete from plan_table;
ORACLE> explain plan for
Select * from BIG where bname like 'BET%';
SQL> select * from table (dbms_xplan.display)
3. Using the Optimization type of FIRST_ROWS then ALL_ROWS do:
ORACLE> delete from plan_table;

ORACLE> explain plan for Select MAX(BIGNO) from big;
SELECT * FROM TABLE (dbms_xplan.display
4. Using the Optimization type of ALL_ROWS then FIRST_ROWS do:
ORACLE> delete from plan_table;
ORACLE> explain plan for
Select * from big where bigno IN (500, 99500);
SELECT * FROM TABLE (dbms_xplan.display

-
5. Using the Optimization type of ALL_ROWS and FIRST_ROWS_1000 do:
ORACLE> delete from plan_table;

ORACLE> SELECT * FROM BIG WHERE BIGNO = 1000 AND
BNAME = 'ONE THOUSAND';

Lab 9- continued

Given non-unique indexes on all columns of the BIG2 table plus a concatenated non-unique index on the BNAME and BIGNO columns of the BIG2 table, use EXPLAIN PLAN to determine how ORACLE10G would execute the following queries:

6. Using the Optimization type of ALL_ROWS and FIRST_ROWS do:
ORACLE> delete from plan_table;

ORACLE> Select bigno from big2 where bname = 'ONE THOUSAND';
7. Using the Optimization type of ALL_ROWS do:
ORACLE> delete from plan_table;

ORACLE> select bigno, bname from big2
where lower(BNAME) = 'one thousand';
8. Using the Optimization type of ALL_ROWS do:
ORACLE> delete from plan_table;

ORACLE> select bigno, bname
from big2
where bname = ' ONE THOUSAND' and bigno = 1;
9. Using the Optimization type of ALL_ROWS and FIRST_ROWS do:
ORACLE> delete from plan_table;

ORACLE> select big.bigno, big2.bname
from big, big2
where big.bigno = ABS(big2.bigno + 0);

(THE DRIVING TABLE IS: _____)

10. Using the Optimization type of ALL_ROWS do:

```
ORACLE> delete from plan_table;
```

```
ORACLE> select big.bigno, big2.bnam
         from big, big2
         where ABS(big.bigno) + 0 = big2.bigno;
```

(THE DRIVING TABLE IS: _____)


11. The following exercise will have you alter you session setting the optimizer to the type of ALL_ROWS and than perform a series of EXPLAIN PLANs for the same SQL statement except for the value being searched. The main purpose of this is to demonstrate when it is more efficient for ORACLE to perform a full table scan versus using an Index on the BIGNO column.

For each of these queries identify the

- | | Cost | Table Access | %Rows Returned |
|----|---------|-------------------------|----------------|
| 1. | ORACLE> | Use the ALL_ROWS Hint; | |
| | ORACLE> | delete from plan_table; | |
| 2. | ORACLE> | SELECT * FROM BIG | |
| | | WHERE BIGNO > 2999999; | |
| 3. | ORACLE> | SELECT * FROM BIG | |
| | | WHERE BIGNO > 2995000; | |
| | ORACLE> | delete from plan_table; | |
| 4. | ORACLE> | SELECT * FROM BIG | |
| | | WHERE BIGNO > 2980000; | |
| | ORACLE> | delete from plan_table; | |
| 5. | ORACLE> | SELECT * FROM BIG | |
| | | WHERE BIGNO > 2971000; | |
| | ORACLE> | delete from plan_table; | |
| 6. | ORACLE> | SELECT * FROM BIG | |
| | | WHERE BIGNO > 2960161; | |
| | ORACLE> | delete from plan_table; | |

-
7. ORACLE> SELECT * FROM BIG
WHERE BIGNO > ?????;
(Where ????? identifies the point at which ORACLE
uses a FULL TABLE SCAN)

PERFORMANCE DIAGNOSTIC TOOLS



- Explain Plan

The actual SQL statement IS NOT executed. Instead, a plan tree is inserted into the PLAN_TABLE. You must run a query on the PLAN_TABLE

- SQL Trace Facility

SQL statement is executed. Use this facility for very long-running queries. This facility records statistics for each SQL statement in a single session, storing statistics in a file with an extension of .trc (Trace Dump File). These statistics are to the nearest second for the PARSE, EXECUTE, and FETCH phase of each SQL statement.

- TKPROF (Trace Kernel Profile)

Submit a trace dump file (.trc) as input to the TKPROF facility. this facility formats a report, and stores it on the disk with an extension of .prf unless you specify your own extension.

SQL Trace Utility



How is the SQL statement performing?

The SQL Trace Utility provides performance information on each SQL statement. It tracks the following statistics:

- - Parse, execute, fetch counts
- - CPU Times & elapsed times
- - Physical & logical reads
- - Number of rows processed

SQL Trace utility is enabled by using INIT.ORA parameters:

- TIMED_STATISTICS=TRUE turns on CPU timings to the nearest second
- TRACE_ENABLED=TRUE to use dump files
- MAX_DUMP_FILE_SIZE=no_blocks file output for trace
- USER_DUMP_DEST=default_directory for dumps
- SQL_TRACE=FALSE (default)
=TRUE (SQL trace facility for every SQL statement)

To activate SQL trace for a single session, enter the following command:

```
ORACLE> ALTER SESSION SET SQL_TRACE = TRUE;
```

```
ORACLE> ALTER SESSION SET SQL_TRACE=FALSE;
```

TKPROF (Trace Kernel Profile)



Produce a readable report

TKPROF translates a trace file into readable format, and can also generate Explain Plan output. Typing TKPROF by itself will give you a on-line help list. This command is executed from a UNIX or DOS prompt.

```
% or c:\orant\bin> TKPROF
```

TKPROF SYNTAX:

```
%TKPROF trace_file output_file [ SORT= sort option1...] print = n  
EXPLAIN=userid/password
```

- trace_file is the input trace file
- output_file is the output file TKPROF has formatted
- SORT sorts the statistics in the output file in sorted order
Some sort options are
 - PRSCNT Count of times parsed
 - PRSCPU Count time during parsing
 - PRSPHR Physical reads during parsing
 - EXECPUT CPU time during execution
 - EXEPHR Physical reads during execution
 - FCHPHR Physical reads during fetch
- PRINT=n output the first n analyzed statements
- EXPLAIN run the EXPLAIN PLAN statement on all of the SQL statements in the trace. Create a Plan table and delete it when finished.

USING TKPROF

1. To use TKPROF, ensure that your session has the SQL Trace Utility set to true:

```
ORACLE> ALTER SESSION SET SQL_TRACE TRUE;
```

2. Identify where the .trc file is. You can identify the user dump file and where it was sent to and if SQL TRACE is TRUE by executing the following script:

```
ORACLE> select name,value from V$PARAMETER
        where NAME = 'user_dump_dest';
```

NAME	VALUE
timed_statistics	TRUE
user_dump_dest	/u01/app/oracle/admin/orcl/udump

3. Execute a sql script which will generate the Trace file with the appropriate statistics.

```
ORACLE> SELECT * FROM BIG
        WHERE BIGNO = 5000;
```

4. ALTER SESSION SET SQL_TRACE=FALSE;

5. Go to the user_dump_dest and identify the .trc file by looking at the date and time it was created. You can use File Manager for this if necessary.

6. Execute TKPROF on the appropriate .trc file.

UNIX

```
%tkprof %ORACLE_HOME%\rdbms\admin\log\ora_1914.trc mytrace
```

```
C:\oracle\product\11.2.0\db_1\BIN\tkprof80 c:\oracle\rdbms\trace\ora_1914.trc mytrace
```

7. Browse the mytrace.prf file created from tkprof in your "CURRENT" directory.

```
c:\oracle\product\11.2.0\db_1\BIN> edit mytrace.prf
```

TKPROF COLUMN DEFINITIONS

TKPROF COLUMNS

CALL	COUNT	CPU	ELAPSED	DISK	QUERY	CURRENT	ROWS

CALL	-	Parse, Execute, or Fetch					
COUNT	-	Number of times a statement was parsed, executed, or fetched. Multiple data blocks may be read by each fetch call.					
CPU	-	Total CPU time in seconds for all parse, execute, or fetch calls for the statement					
ELAPSED	-	Total elapsed time in seconds for all parse, execute, or fetch calls for the statement					
DISK	-	Total number of data blocks physically read from the data files on disk for all parse, execute, or fetch calls					
QUERY	-	Total number of buffers retrieved in consistent mode for all parse, execute, or fetch calls					
CURRENT	-	Total number of buffers retrieved in current mode. Buffers are often retrieved in current mode for insert, update, and delete statements					
ROWS	-	Total number of rows processed by the SQL statement. The total does not include rows processed by subqueries.					

NOTE: For select statements, the number of rows returned appears for the fetch step. For delete, insert, and update statements, the number of rows processed appears for the execute step.

Execute phase is concerned with INSERTS, UPDATES & DELETES

Fetch phase is concerned with SELECTs

SAMPLE TKPROF REPORT

C:\\$ORACLE_HOME\ADMIN\ORCL\TRACE> EDIT mytrace.prf

```
select *  
from staff where id > 50
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.09	0.23	6	25	1	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	3	0.01	0.01	1	3	3	40
<hr/>							
total	5	0.10	0.24	7	28	4	40

Misses in library cache during parse: 1
Optimizer hint: CHOOSE
Parsing user id: 5

OVERALL TOTALS FOR ALL NON-RECURSIVE STATEMENTS

call	count	cpu	elapsed	disk	query	current	rows
Parse	3	0.09	0.23	6	25	1	0
Execute	4	0.00	0.14	0	6	0	0
Fetch	5	0.01	0.01	1	3	3	48

Misses in library cache during parse: 1
Misses in library cache during execute: 1

4 user SQL statements in session.
5 internal SQL statements in session.
10 SQL statements in session.

Processed trace file: ora_1941.trc

1 session in tracefile.
4 user SQL statements in tracefile.
6 internal SQL statements in tracefile.
10 SQL statements in tracefile.
9 unique SQL statements in tracefile.
109 lines in trace file.

LAB 10 - USING SQL TRACE AND TKPROF



Using SQL Trace Utility and TKPROF, compare the following queries to see which one appears to be the fastest:

1. ORACLE> SELECT BIG.BNAME, BIG2.BIGNO
 FROM BIG, BIG2
 WHERE BIG.BIGNO = BIG2.BIGNO AND
 BIG2.BNAME = 'ONE THOUSAND';

2. ORACLE> SELECT BIG.BNAME
 FROM BIG
 WHERE BIGNO IN (SELECT BIGNO FROM
 BIG2
 WHERE BNAME = 'ONE THOUSAND');

3. ORACLE> SELECT BIG.BNAME
 FROM BIG
 WHERE 0 < (SELECT COUNT(*)
 FROM BIG2
 WHERE BIG2.BNAME = 'ONE THOUSAND'
 AND BIG.BIGNO = BIG2.BIGNO);

LAB 10 - USING SQL TRACE AND TKPROF (CONTINUED)

4. Fill in the formatted report from TKPROF for the first SQL statement:

```
c:\tkprof c:\orant\rdms80\admin\TRACE\ora_____.trc lab84
explain=userid/password
```

lab84.prf

```
select big.bname, big2.bigno
from big, big2
where big.bigno = big2.bigno
and big2.bname = 'ONE THOUSAND';
```

TKPROF COLUMNS

CALL	COUNT	CPU	ELAPSED	DISK	QUERY	CURRENT	ROWS
Parse							
Execute							
Fetch							

Misses in library cache during parse: _____

Parsing user id: _____userid _____

ROWS	EXECUTION PLAN

LAB 10 - USING SQL TRACE AND TKPROF (CONTINUED)

5. Fill in the formatted report from TKPROF for the first SQL statement:

c:\tkprof c:\oracle\rdbms\admin\TRACE\ora_____.trc lab84

explain=**userid/password**

lab84.prf

```
select big.bname,  
from big  
where BIGNO IN (SELECT BIGNO  
                FROM BIG2  
                WHERE BNAME = 'ONE THOUSAND');
```

TKPROF COLUMNS

CALL	COUNT	CPU	ELAPSED	DISK	QUERY	CURRENT	ROWS
Parse							
Execute							
Fetch							

Misses in library cache during parse: _____

Parsing user id: _____userid _____

ROWS	EXECUTION PLAN

LAB 10 - USING SQL TRACE AND TKPROF (CONTINUED)

6. Fill in the formatted report from TKPROF for the first SQL statement:

```
c:\tkprof c:\$ORACLE_HOME\rdbms\admin\orcl\udump\ora_____.trc lab84
explain=userid/password
```

lab84.prf

```
select big.bname
from big
where 0 < (SELECT COUNT(*)
          FROM BIG2
 WHERE big2.bname = 'ONE THOUSAND'
       AND BIG.BIGNO = BIG2.BIGNO);
```

TKPROF COLUMNS

CALL	COUNT	CPU	ELAPSED	DISK	QUERY	CURRENT	ROWS
Parse							
Execute							
Fetch							

Misses in library cache during parse: _____
Parsing user id: _____userid _____

ROWS	EXECUTION PLAN

This page left blank intentionally